

Rewa Parashar

rewaparashar5@gmail.com | +91 9975059011 | [GitHub](#) | [LinkedIn](#) | Pune, Maharashtra

SUMMARY

Data Engineer with 1.5+ years of hands-on experience building and optimizing ETL/ELT pipelines and data quality solutions on the Azure + Databricks stack (Spark, Delta Lake, LakeFlow, ADF, ADLS). Contributed to a 20% pipeline runtime improvement and 99.9% processing reliability at Parkar Global Technologies. Experienced in data governance, monitoring, and collaborating with business and IT stakeholders to deliver reliable, scalable data platforms. Databricks Certified Data Engineer Professional.

WORK EXPERIENCE

Parkar Global Technologies

January 2024 – Present

Software Engineer I — Data Engineering

Pune, India

- Built a **Spark Structured Streaming + LakeFlow** declarative pipeline supporting incremental ingestion, Delta Lake ACID transactions, and schema-aware transformations in Databricks, processing high-volume raw data across multiple downstream analytics consumers.
- Improved **ETL/ELT pipeline runtime by 20%** by optimizing Spark transformations (predicate pushdown, broadcast joins, partition pruning) and tuning cluster auto-scaling configurations, reducing average job latency from 4.2 hours to 2.9 hours.
- Implemented continuous **data quality monitoring and integrity checks** within pipelines, enabling early detection of data anomalies and reducing downstream reporting errors; documented DQ frameworks for team reusability and knowledge transfer.
- Reduced manual data engineering effort by **40%** through reusable declarative pipeline templates, cutting new pipeline onboarding time from 3 days to under 1 day; participated in **Agile/Scrum** sprints and maintained solution documentation for stakeholder transparency.
- Contributed to increasing data processing reliability to **99.9% success rate** leveraging declarative orchestration with automated retries and self-healing recovery, eliminating weekly manual incident interventions.
- Worked alongside senior engineers to implement **RBAC/ABAC data governance policies** with PII masking, metadata management, and secure data pipelining across regulated datasets, supporting audit-ready compliance requirements.
- Collaborated with business and IT stakeholders to translate data requirements into pipeline designs, conducted **solution validation testing** against functional and non-functional requirements, and partnered with the architecture team on data storage design decisions.

CERTIFICATIONS

- Databricks Certified Data Engineer Professional** - Databricks
- Databricks Certified Data Engineer Associate** - Databricks
- Microsoft Azure Fundamentals (AZ-900)** - Microsoft

SKILLS

Distributed Systems: Apache Spark | Delta Lake | LakeFlow DSP (Delta Live Tables) | Spark Structured Streaming
Data Engineering: ETL/ELT | Data Modelling | Schema Evolution | Workflow Orchestration | Batch & Streaming | Data Profiling | Data Quality Monitoring
Programming: Python | PySpark | SQL | Spark SQL
Cloud & Infrastructure: Azure (ADLS, ADF, Azure SQL Database, Databricks) | AWS S3
Data Governance: RBAC | ABAC | PII Masking | Metadata Management | Secure Data Pipelining
Visualization: Power BI | Databricks Dashboard | Streamlit
DevOps & Agile: Git | GitHub | GitLab | Agile/Scrum | Kanban

PROJECTS

Financial Banking Data Pipeline & Analytics Platform *(Personal Project)*

- Designed and built a **Spark + LakeFlow** pipeline end-to-end to deepen understanding of regulated data ingestion patterns — modeled after real banking compliance and data governance requirements, including metadata management and access control.
- Implemented **data quality checks and anomaly detection** using a Random Forest classifier on synthetic transactional data, achieving 94% detection accuracy on a held-out test set; documented the DQ framework for reproducibility.
- Visualized pipeline output via **Power BI** dashboards surfacing key transaction KPIs, and built a Streamlit-based what-if interface to simulate how analysts might investigate flagged anomalies without engineering intervention.

F1 Insights Analytics Pipeline *(Personal Project)*

- Developed an end-to-end **PySpark analytics pipeline** processing 70+ years of Formula 1 race data (~2M records) covering results, lap times, pit stops, and constructor standings.
- Built transformation workflows generating aggregated metrics on driver performance, qualifying efficiency, and pit strategy effectiveness, cutting query time by **65%** vs. raw CSV access.
- Conducted statistical correlation analysis (tire choice vs. grid position vs. race outcome) and visualized findings in interactive dashboards.

EDUCATION

Bachelor of Engineering – Information Technology

Savitribai Phule Pune University | GPA: 9.1 / 10.0

June 2020 – June 2024